

---

# Two Machine Learning Tasks on Spark platform

---

Anonymous Author(s)

Affiliation

Address

email

## Abstract

1 Facing the era of data explosion in which massive data are created, stored and  
2 analyzed, Spark comes as a new framework for distributed computing. Spark pro-  
3 vides native functional APIs that support multiple programming languages. In this  
4 paper, we elaborate on two machine learning experiments, facial keypoints detec-  
5 tion and large scale video classification, conducted on Spark platform. Different  
6 models are deployed and compared based on performance. Through practical ex-  
7 perience of programming on Spark, we get familiar with concepts like resilient  
8 distributed datasets (RDD), etc. and detailed mechanism of distributed comput-  
9 ing.

## 1 Detect the location of keypoints on face images

### 1.1 Introduction

12 Face detection is a classical and essential problem in computer vision and pattern recognition. It  
13 tackles gender recognition, emotion recognition and other real-life tasks that commonly occur  
14 in facial expression analysis, object tracking, etc. The challenges mainly come from the variety  
15 of illumination, rotation and pose. In this part, we propose a facial keypoints detection question  
16 retrieved from [www.kaggle.com](http://www.kaggle.com). We try to predict the locations of points including eyes, nose  
17 and mouth from  $96 \times 96$  pixels of an image. We see it as a regression problem so as to fit a  
18 regression model on the pixel features. For simplicity, we call off-the-shelf interfaces in Spark  
19 library to complete this work.

### 1.2 Data preprocessing

21 Noticing there are many data with value NA in the location columns, we fill them with the average  
22 of the corresponding column. This may not be reasonable, but if we drop all instances with missing  
23 data, only 30% of training instances will be left. Besides, we scale the range of pixels to  $[0, 1]$  by  
24 zero mean normalization. Also, input files are reorganized so that they can be transformed to RDD  
25 properly.

26 As for feature extraction, initially we sample 3,000 dimensions out of 9,216 but the computer fails  
27 to give results and throws memory error. Therefore, we decide to use 512-dimension pixel features  
28 from random sampling at last. In addition, we deploy PCA for dimension reduction. Note that SVD  
29 upon such a matrix is infeasible so we implement PCA over the sampled features.

### 1.3 Models

31 In this section we briefly illustrate some regression models employed in the experiments later.

32 **Isotonic regression** Isotonic regression is a generalized form of linear regression that fits a non-  
33 decreasing function to data. The free-form property of isotonic regression makes it outperform linear  
34 regression sometimes because it fits the data points in a more elastic way.

35 **Tree-based models** Tree-based regressions are all in relation to decision trees to some extent.  
36 In particular, gradient boosted tree regressor uses gradient boosting iteration to fit a decision tree.  
37 However, it can hardly be parallelized due to the sequential nature of boosting. Random forest  
38 regressor outputs the mean prediction of individual trees, which avoids overfitting habits of decision  
39 trees.

## 40 1.4 Experiments

41 Since we do not have the ground truth of the test data, we randomly split the training set by 9 : 1 for  
42 training and testing respectively. The metric we use to evaluate the results is root of mean squared  
43 error (RMSE), i.e. the root of average deviations between the true location coordinate and the  
44 prediction.

45 We load the data into RDD and transform features into vectors so they can be valid arguments of  
46 models in Spark. Then we call different modules from the library `pyspark.ml.regression`. For  
47 each model, we first construct a model instance and fit it over the training data. After that we  
48 make predictions by calling `transform()` method. The predictions are evaluated by evaluator in  
49 `pyspark.ml.evaluation`.

50 To compare models, we take “left\_eye\_center\_x” as example and show the results in the following  
table.

Table 1: RMSE of predictions

Model	RMSE
Linear regression	3.56
Decision tree	3.80
Random forest	3.35
Gradient boosted tree	3.62
Isotonic regression	3.54

51

## 52 1.5 Discussion

53 In Table 1, random forest regressor gives the best result. It outperforms decision tree just as expected,  
54 indicating overfitting may occur in the latter case. Gradient boosted tree regressor is said to be able  
55 to handle data of mixed type, which cannot be ratified in this experiment. Linear regression, as a  
56 basic method in regression, yields intermediate performance among all. Here we observe a subtle  
57 but significant gap between isotonic regression and linear regression. If we look into the pairs of the  
58 true value and prediction, we can see in some cases the error is relatively small ( $< 1$ ). This implies  
59 there are cases in which our regression models do not work well.

60 Considering loss of information due to feature dimension sampling and training instance sampling,  
61 we can expect better performance with more efficiency of data usage. On the other hand, we can  
62 introduce features which are more sophisticated, such as pretrained CNN features. Meanwhile,  
63 efforts can also be made on a technical level, i.e. utilizing neural network models instead of basic  
64 regression models.

## 65 2 Large-scale video classification

### 66 2.1 Introduction

67 Video is one of the most common forms of multimedia nowadays. Videos are posted, transferred  
68 and viewed throughout the Internet and users' terminal devices. Video classification is a technique  
69 that attributes videos of similar theme or content to one category. It is widely applied in content  
70 search and recommendation of online videos.

71 For efficient classification, multiple features should be taken into consideration. Static appearance in-  
72 formation and acoustic channels, motion clues are mutually complementary in representing a video.  
73 As for the former, CNN based representations are usually selected as static features. Motion features  
74 extends frame-based local features into 3D space. One can locate densely sampled frame patches to  
75 generate dense trajectories.

76 We work on Fudan-Columbia Video Dataset (FCVID) which contains over 90k Internet videos with  
77 239 manually annotated categories. We utilize classification methods including naïve Bayes, multi-  
78 class logistic regression, etc. and compare the prediction accuracy. We also study inter-class rela-  
79 tionships which are said to boost classification performance but also cause confusion. For example,  
80 class "WeddingReception" and "WeddingDance" may be correlated.

### 81 2.2 Features

82 **CNN** CNN is proved successful on image feature extraction in which pixels are filtered by feature  
83 map. However, its performance in video classification meets the bottleneck caused by complexity  
84 and scale of videos. Works have been done to extend the CNN to exploit more information, such as  
85 spatial-temporal space (Ji *et al.*, 2010), two-stream CNN (Simonyan *et al.*, 2014) and some advanced  
86 feature encoding strategies (Xu *et al.*, 2015).

87 **HOG** Histogram of oriented gradients (HOG) is a feature descriptor that counts occurrences of  
88 gradient orientation in localized portions of a video frame. It is computed on a dense grid of uni-  
89 formly spaced cells and uses overlapping local contrast normalization to improve accuracy. The first  
90 step of descriptor calculation is computing gradients. The next step is to create the cell histograms.  
91 Every pixel within the cell delivers a weighted vote for an orientation-based histogram channel based  
92 on the values in the gradient computation.

93 **SIFT** Scale-invariant feature transform (SIFT) is a descriptor to detect local features in images or  
94 video frames. It is able to extract features in a way that reduces contribution of errors brought by  
95 image scale, noise and illumination. SIFT keypoints of objects are first generalized from a set of  
96 reference images. An object is recognized in a new image by comparing each feature from the new  
97 image to the forementioned candidates and finding cluster based on Euclidean distance of feature  
98 vectors. The clusters are determined by Hough transform.

### 99 2.3 Experiments

100 We obtain features from [http://bigvid.fudan.edu.cn/data/fcvid/FCVID\\_Feature/](http://bigvid.fudan.edu.cn/data/fcvid/FCVID_Feature/). For  
101 each feature, we sample 512 dimensions. Likewise, we randomly sample 5,000 training instances  
102 among all 91,223. The ratio of training and testing data is 9 : 1. The evaluation metric is accuracy.  
103 Our experiments can be divided into the following parts:

104 We first apply different classifiers provided by `pyspark.ml.classification` library on the  
105 training set. We examine the performance of models by comparing accuracy of predic-  
106 tions they make based on the testing set, using `MulticlassClassificationEvaluator` in  
107 `pyspark.ml.evaluation`.

108 Then we check how features work complementarily. Features are deployed both individually and  
109 together to see whether feature fusion guarantees improvement in performance. The fusion method  
110 we use is simple concatenation of feature vectors.

111 Next, we change the scale of the training set in order to observe the effect on performance due  
112 to sample size. The number of instances ranges from 450 to 2,250. This part of experiment is  
113 implemented over multiclass logistic regression model.

114 Finally, we study inter-class relationships to mining knowledge sharing in classes.

## 115 2.4 Results

116 We display the comparison of classifiers in the table below.

Table 2: Accuracy of predictions (training intance: 1800)

Model	Accuracy
Logistic regression	0.31
Decision tree	0.07
Naïve Bayes	0.33
Random forest	0.06
Multilayer perceptron	0.17

117 In terms of features, we test different feature combinations.

Table 3: Features and accuracy (model: logistic)

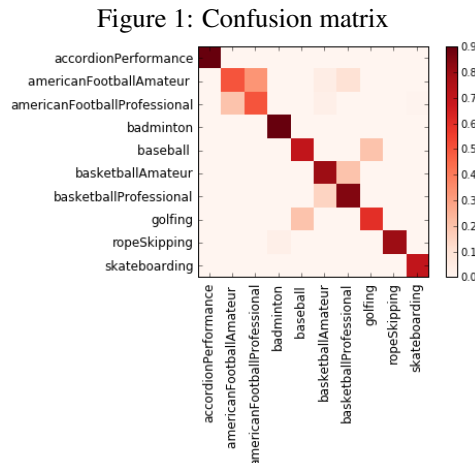
Feature(s)	Accuracy
CNN	0.29
HOG	0.27
SIFT	0.26
CNN+HOG	0.30
CNN+SIFT	0.32

118 As training instances gradually increase , the accuracy keeps going up as Table 3 shows.

Table 4: Sample size and accuracy (model: logistic)

Sample size	Accuracy
450	0.15
900	0.27
1,350	0.30
1,800	0.31
2,250	0.37

119 Inter-class relationships are visualized with the confusion matrix below.



## 2.5 Discussion

It should be mentioned at the very beginning that two very low values in Table 2 result from the lack of training instances. Once we train these two tree-based models on a larger set, we can observe improvement of accuracy immediately.

Two relatively simple models, naïve Bayes and multiclass logistic regression work effectively even with a small training set. The MLP and tree-based methods need more instances to learn a better model. Moreover, the latter may incur overfitting because we do not go through the pruning process.

About feature fusion, it is verified that feature combinations are able to promote the prediction accuracy. In fact, HOG supplements information in motion because CNN features are usually considered as static presentations of an image. Besides, SIFT carries audio descriptors that help categorize video semantics. Therefore, the three features we employ in this experiment contain static, motion and audio information respectively and are complementary to each other.

At last, we comment on Figure 1. The colored spots which do not lie on the diagonal indicate that the corresponding two classes are correlated while may bring confusion in classification. The knowledge sharing property is identified by matrix  $\Omega$  (Jiang *et al.*, 2015).

## References

- [1] Y. Jiang, Z. Wu, J. Wang, X. Xue & S. Chang (2015) Exploiting feature and class relationships in video categorization with regularized deep neural networks. In *ArXiv preprint*.
- [2] S. Ji, W. Xu, M. Yang & K. Yu (2010) 3d convolutional neural networks for human action recognition. In *ICML*.
- [3] K. Simonyan & A. Zisserman (2014) Two-stream convolutional networks for action recognition in videos. In *NIPS*.